

# **Predictive Credit Score Classification Model**

## **Using Machine Learning and Neural Networks**

### **Project 1**

**Urvashi Dube**

**NUID: 002752858**

[dube.u@northeastern.edu](mailto:dube.u@northeastern.edu)

**Submission Date:** 29<sup>th</sup> November, 2023

## Objective

The primary objective of this project is to establish a robust predictive credit score classification model by employing a combination of machine learning techniques and neural networks. A diverse range of customer features, including but not limited to annual income, age, occupation, and payment behavior, will be analyzed by the model to categorize individuals into credit score classes such as "Good," "Bad," or "Standard." By exploring a comprehensive dataset encompassing various financial parameters, the aim is to develop an accurate and reliable tool for assessing creditworthiness. This tool will provide valuable insights for risk management and decision-making in the financial domain and help companies take data driven decisions.

## Stakeholders

The analysis is anticipated to draw attention from various parties, including banks, financial institutions, credit card companies, and individuals seeking loans or credit. The implementation of the predictive credit score classification model is expected to provide a reliable means of assessing the creditworthiness of potential borrowers. Through the consideration of various factors such as annual income, age, occupation, and payment behavior, individuals are categorized into credit score classes, including "Good," "Bad," or "Standard." A thorough analysis of a diverse dataset containing financial information is done using machine learning and neural networks. This tool is deemed a valuable resource for stakeholders, aiding them in making well-informed decisions in areas such as risk management and financial planning. Additionally, the model can be utilized by individuals to check their own credit scores, enabling them to make informed choices for significant financial commitments, such as purchasing homes or vehicles.

## Importing data

Initially, various python libraries were imported, these include Pandas for data handling, NumPy for numerical operations, Seaborn and Matplotlib for data visualization, and Plotly for interactive plotting. The scikit-learn library is employed for machine learning tasks, featuring modules such as `DecisionTreeClassifier`, `RandomForestClassifier`, and `SGDClassifier` for classification purposes. Additionally, XGBoost, a powerful gradient boosting library, is imported. The code also integrates imputation techniques such as `KNNImputer` from scikit-learn to handle missing data and `SMOTE` from `imbalanced-learn` for oversampling to address class imbalance. Furthermore, statistical functions from SciPy and metrics from scikit-learn are included for model evaluation.

The train and test datasets are read from CSV files, with a limited number of rows (20000 train rows and 10000 test rows) imported due to hardware constraints. This initial data exploration and preparation phase sets the foundation for subsequent stages of the project, providing a comprehensive toolkit for data analysis and machine learning model development. The screenshot of first 5 rows of train data is shown in Fig. 1. and test data is shown in Fig. 2.

	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts	...	Credit_Mix	Outstanding
0	0x1602	CUS_0xd40	January	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	3	...	...	...
1	0x1603	CUS_0xd40	February	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	NaN	3	...	Good	...
2	0x1604	CUS_0xd40	March	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	NaN	3	...	Good	...
3	0x1605	CUS_0xd40	April	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	NaN	3	...	Good	...
4	0x1606	CUS_0xd40	May	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	3	...	Good	...

Fig. 1. Glimpse of Train Data

	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts	...	Num_Credit_Inquiries
0	0x160a	CUS_0xd40	September	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	3	...	2022.0
1	0x160b	CUS_0xd40	October	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	1824.843333	3	...	4.0
2	0x160c	CUS_0xd40	November	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	1824.843333	3	...	4.0
3	0x160d	CUS_0xd40	December	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	NaN	3	...	4.0
4	0x1616	CUS_0x21b1	September	Rick Rothackerj	28	004-07-5839	...	34847.84	3037.986667	2	...	5.0

Fig. 2. Glimpse of Test Data

## Preprocessing Data

Data preprocessing is first conducted on a dataset named 'train\_df'. The initial steps involve examining the dataset's shape, information, and summary statistics to gain insights into its structure. The dataset consists of 28 columns, with some columns containing null values that needed to be addressed. The value counts of specific categorical columns can be seen in Fig. 3.

```

Occupation
Mechanic      1292
Lawyer        1286
Doctor        1235
Journalist    1197
Developer     1192
Media_Manager 1188
Teacher       1170
Entrepreneur  1159
Musician      1154
Architect     1140
Writer        1126
Scientist     1118
Engineer      1112
Manager       1096
Accountant    1093
Name: count, dtype: int64

Credit_Mix
Standard      7180
Good          5003
Bad           3810
Name: count, dtype: int64

Payment_of_Min_Amount
Yes           9677
No            6633
NM            2235
Name: count, dtype: int64

Month
February      2332
January        2331
July           2327
March          2321
May            2314
June           2312
August         2311
April          2297
Name: count, dtype: int64

0      High_spent_Small_value_payments
1      Low_spent_Large_value_payments
2      Low_spent_Medium_value_payments
3      Low_spent_Small_value_payments
4      High_spent_Medium_value_payments
Name: Payment_Behaviour, dtype: object

0      4
1      3
2      2
3      1
4      5
Name: Payment_Behaviour, dtype: int64

```


Fig. 3. Value counts of category data

The following preprocessing techniques were applied on the data:

- Unnecessary columns, including "ID," "Name," and "SSN," are dropped in the data preparation phase.
- Data cleaning is applied to columns like "Credit\_Mix" and "Changed\_Credit\_Limit," involving the replacement of specific values and addressing missing data.
- Categorical columns are explored, and rows containing the value "!!@9#%8" in the "Payment\_Behaviour" column are identified and removed.
- Investigation and potential transformation of count-based columns, such as "Num\_of\_Loan" and "Num\_of\_Delayed\_Payment," are conducted.
- The "Age" column undergoes processing to extract meaningful information, and missing values are appropriately handled.
- Similar operations are performed on columns like "Annual\_Income," "Num\_of\_Loan," and "Outstanding\_Debt" to ensure data quality.
- The "Credit\_History\_Age" column is processed to extract years and months, resulting in the creation of new columns, namely "Credit\_Age\_years" and "Credit\_Age\_months."
- Label encoding is applied to the "Payment\_Behaviour" column, and missing values are filled using forward fill.


- Adjustment of the "Num\_Bank\_Accounts" column involves replacing -1 with 0 and rounding to the nearest integer.
- The "Type\_of\_Loan" column is filled with forward fill to handle missing values.
- Outliers and invalid values in various columns, including "Delay\_from\_due\_date," "Credit\_Utilization\_Ratio," and "Total\_EMI\_per\_month," are addressed during the data preparation process.

A heatmap is generated to visualize missing values, highlighting areas that require imputation (can be seen in Fig. 4).



**Fig. 4. Missing Values heatmap before KNN imputer**

The missing data is then imputed using KNNImputer (can be seen in Fig. 5) for numerical columns and IterativeImputer for categorical columns, ensuring the dataset is ready for further analysis.



**Fig. 5. Missing Values heatmap after KNN imputer**

A glimpse of the processed train data can be seen in Fig. 6. without the null values. The null values are imputed using KNN Imputer which could be seen in the heatmap.

```
train_df.head()
```

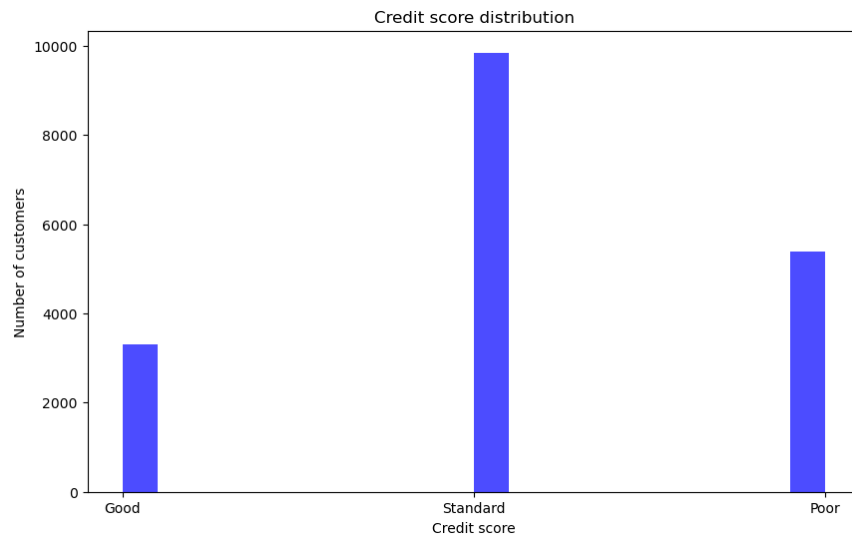
nent_of_Min_Amount	Total_EMI_per_month	Amount_invested_monthly	Payment_Behaviour	Monthly_Balance	Credit_Score	Credit_Age_years	Credit_Age_months
No	49.575	80.415	4	312.494	Good	22.0	1.0
No	49.575	118.280	3	284.629	Good	NaN	NaN
No	49.575	81.700	2	331.210	Good	22.0	3.0
No	49.575	199.458	1	223.451	Good	22.0	4.0
No	49.575	41.420	5	341.489	Good	22.0	5.0

**Fig. 6. Glimpse of train data**

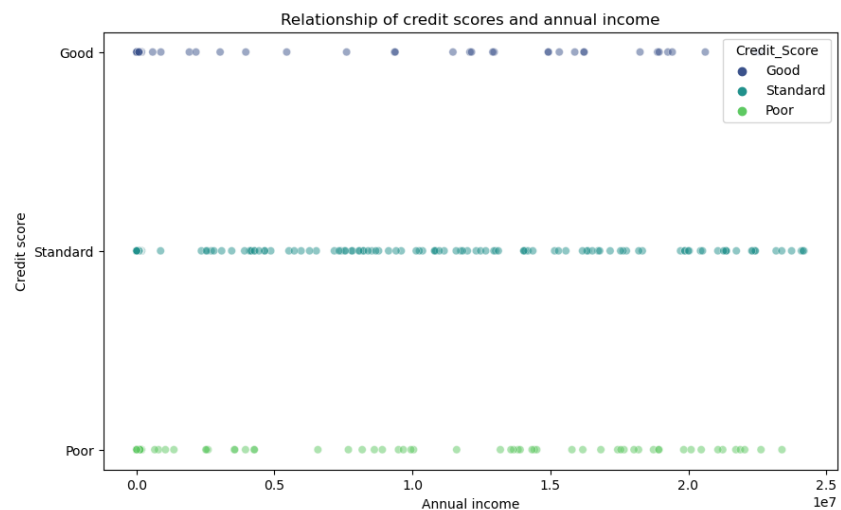
Finally, label encoding is applied to the "Credit\_Mix" column, mapping categorical values to numerical representations.

## EDA Analysis

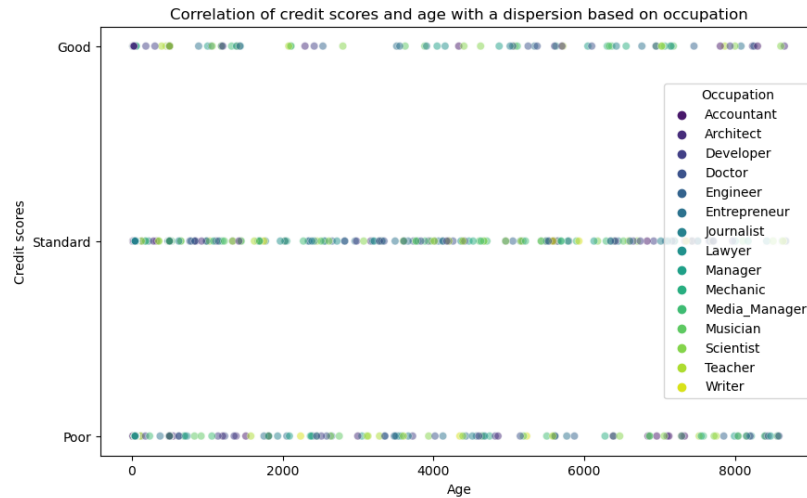
The dataset (`train_df`) is initially explored for missing values, and unique values for each column are printed. The distribution of credit scores (seen in Fig. 7), credit scores against annual income (shown in Fig. 8), and the correlation of credit scores with age based on occupation are also examined (shown in Fig. 9). The distribution of the credit scores is seen in the bar chart. There were more “standard” credit scores in the data.



**Fig. 7. Credit Score Distribution**




**Fig. 8. Credit Score against annual income Distribution**




**Fig. 9. Correlation of credit scores and age with a dispersion based on occupation**

Insights into the distribution of payment behavior types are provided using both a count plot and a pie chart seen in Fig. 10 and Fig. 11. “Lawyer” and “Mechanic” occupation are most seen in the bar chart and least occupation is “Accountant”. In the next scatter graph “Accountant” are proved to have good credit score and “Writer” have Poor credit score. Payment behavior of “1” is seen to be the most in the data.




**Fig. 10. Distribution of Payment Behavior Types**






**Fig. 11. Occupation counts in the data**

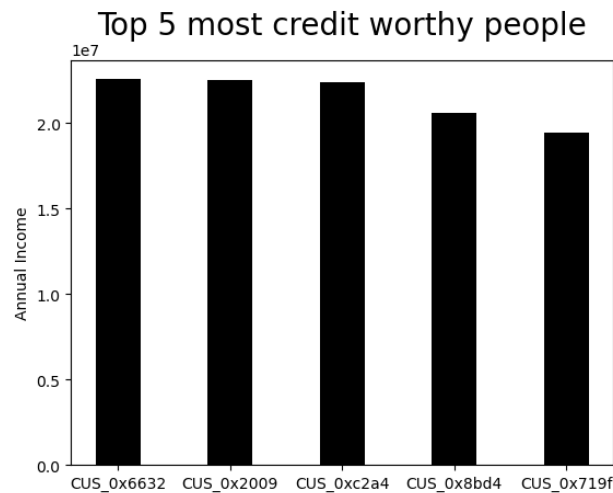
Further visualizations encompass a scatter plot depicting the correlation between credit scores and the number of delays from the due date (seen in Fig. 12), a count plot for a categorical feature related to the payment of the minimum amount (seen in Fig. 13), and a bar chart showcasing the top 5 most credit-worthy individuals based on annual income (seen in Fig. 14). The next scatter plot shows that delays in payments have led to poor credit score in the data.



**Fig. 12. Correlation of credit scores and the number of delay from due date**



**Fig. 13. Count of Payment of Minimum amount in the data**



**Fig. 14. Top 5 most credit worthy people**

The bar chart data clearly shows that the payment of minimum amount is done by maximum number of people. The data from top 5 bar chart graph shows that “CUS\_0x6632” is the most credit worthy.

A heatmap of the correlation matrix for numerical features in the dataset is generated, offering an overview of the relationships between different variables (seen in Fig. 15). The heatmap clearly shows that monthly inhand salary and amount invested monthly are highly correlated. Which is also seen in dealy from due date and outstanding debt. Highest negative correlation is seen in between outstanding debt and monthly balance. This is also seen in outstanding debt and monthly inhand salary where the negative correlation is of -0.26.

## IE 7300: Statistical Learning for Engineering




Fig. 15. Correlation heatmap between columns

A glimpse of the description train data can be seen in Fig. 16.

```
train_df.describe()
```

	Customer_ID	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate
count	18545.000000	18545.000000	18545.000000	18545.000000	18545.000000	18545.000000	18545.000000	18545.000000	18545.000000
mean	1248.172284	3.495282	23.622378	7.044702	1263.763440	2143.353572	6.886384	10.329954	17.128498
std	721.119934	2.293685	32.898839	4.255661	729.922935	1284.181183	15.421818	35.050791	31.637957
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	625.000000	1.000000	11.000000	3.000000	632.000000	1013.000000	3.000000	4.000000	6.000000
50%	1247.000000	3.000000	20.000000	7.000000	1260.000000	2129.000000	6.000000	6.000000	12.000000
75%	1872.000000	6.000000	28.000000	11.000000	1895.000000	3259.000000	8.000000	7.000000	19.000000
max	2499.000000	7.000000	396.000000	14.000000	2681.000000	4401.000000	233.000000	395.000000	385.000000

8 rows × 26 columns

Fig. 16. Description of Train Data

The distribution of numerical features concerning credit scores is visualized using a box plot (can be seen in Fig. 17(a), (b), (c)). Chi-square tests is conducted to assess the independence between

each categorical variable and the target variable (credit score) (a glimpse of the chi squared test can be seen in Fig. 18). The graph clearly shows that there are many outliers in age, num bank accounts, num credit card, interest rate, number of delayed payments, number of credits inquires. The chi squared test shows the following insights:

- Customer\_ID: The chi-square statistic is 24697.31, with a p-value of 0.0 and 4998 degrees of freedom. This indicates a significant association between the Customer\_ID and the outcome variable.
- Month: The chi-square statistic is 24.20, with a p-value of 0.043 and 14 degrees of freedom. This suggests a significant association between the Month and the outcome variable, although the association is weaker compared to Customer\_ID.
- Age: The chi-square statistic is 1844.21, with a p-value close to 0 ( $1.145e-85$ ) and 792 degrees of freedom. This indicates a significant association between Age and the outcome variable.
- Occupation: The chi-square statistic is 97.66, with a p-value close to 0 ( $1.208e-09$ ) and 28 degrees of freedom. This suggests a significant association between Occupation and the outcome variable.
- Annual\_Income: The chi-square statistic is 24857.72, with a p-value of 0.0 and 5362 degrees of freedom. This indicates a significant association between Annual\_Income and the outcome variable.
- Monthly\_Inhand\_Salary: The chi-square statistic is 26349.37, with a p-value of 0.0 and 8802 degrees of freedom. This suggests a significant association between Monthly\_Inhand\_Salary and the outcome variable.
- Num\_Bank\_Accounts: The chi-square statistic is 4288.21, with a p-value of 0.0 and 466 degrees of freedom. This indicates a significant association between Num\_Bank\_Accounts and the outcome variable.

- Num\_Credit\_Card: The chi-square statistic is 5057.28, with a p-value of 0.0 and 790 degrees of freedom. This suggests a significant association between Num\_Credit\_Card and the outcome variable.
- Interest\_Rate: The chi-square statistic is 7094.01, with a p-value of 0.0 and 770 degrees of freedom. This indicates a significant association between Interest\_Rate and the outcome variable.
- Num\_of\_Loan: The chi-square statistic is 2680.41, with a p-value of 0.0 and 86 degrees of freedom. This suggests a significant association between Num\_of\_Loan and the outcome variable.
- Type\_of\_Loan: The chi-square statistic is 14964.27, with a p-value of 0.0 and 3072 degrees of freedom. This indicates a significant association between Type\_of\_Loan and the outcome variable.
- Delay\_from\_due\_date: The chi-square statistic is 4548.88, with a p-value of 0.0 and 134 degrees of freedom. This suggests a significant association between Delay\_from\_due\_date and the outcome variable.
- Num\_of\_Delayed\_Payment: The chi-square statistic is 3931.89, with a p-value of 0.0 and 598 degrees of freedom. This indicates a significant association between Num\_of\_Delayed\_Payment and the outcome variable.

In summary, all the variables tested show a significant association with the outcome variable. However, the strength of the association varies across the variables.

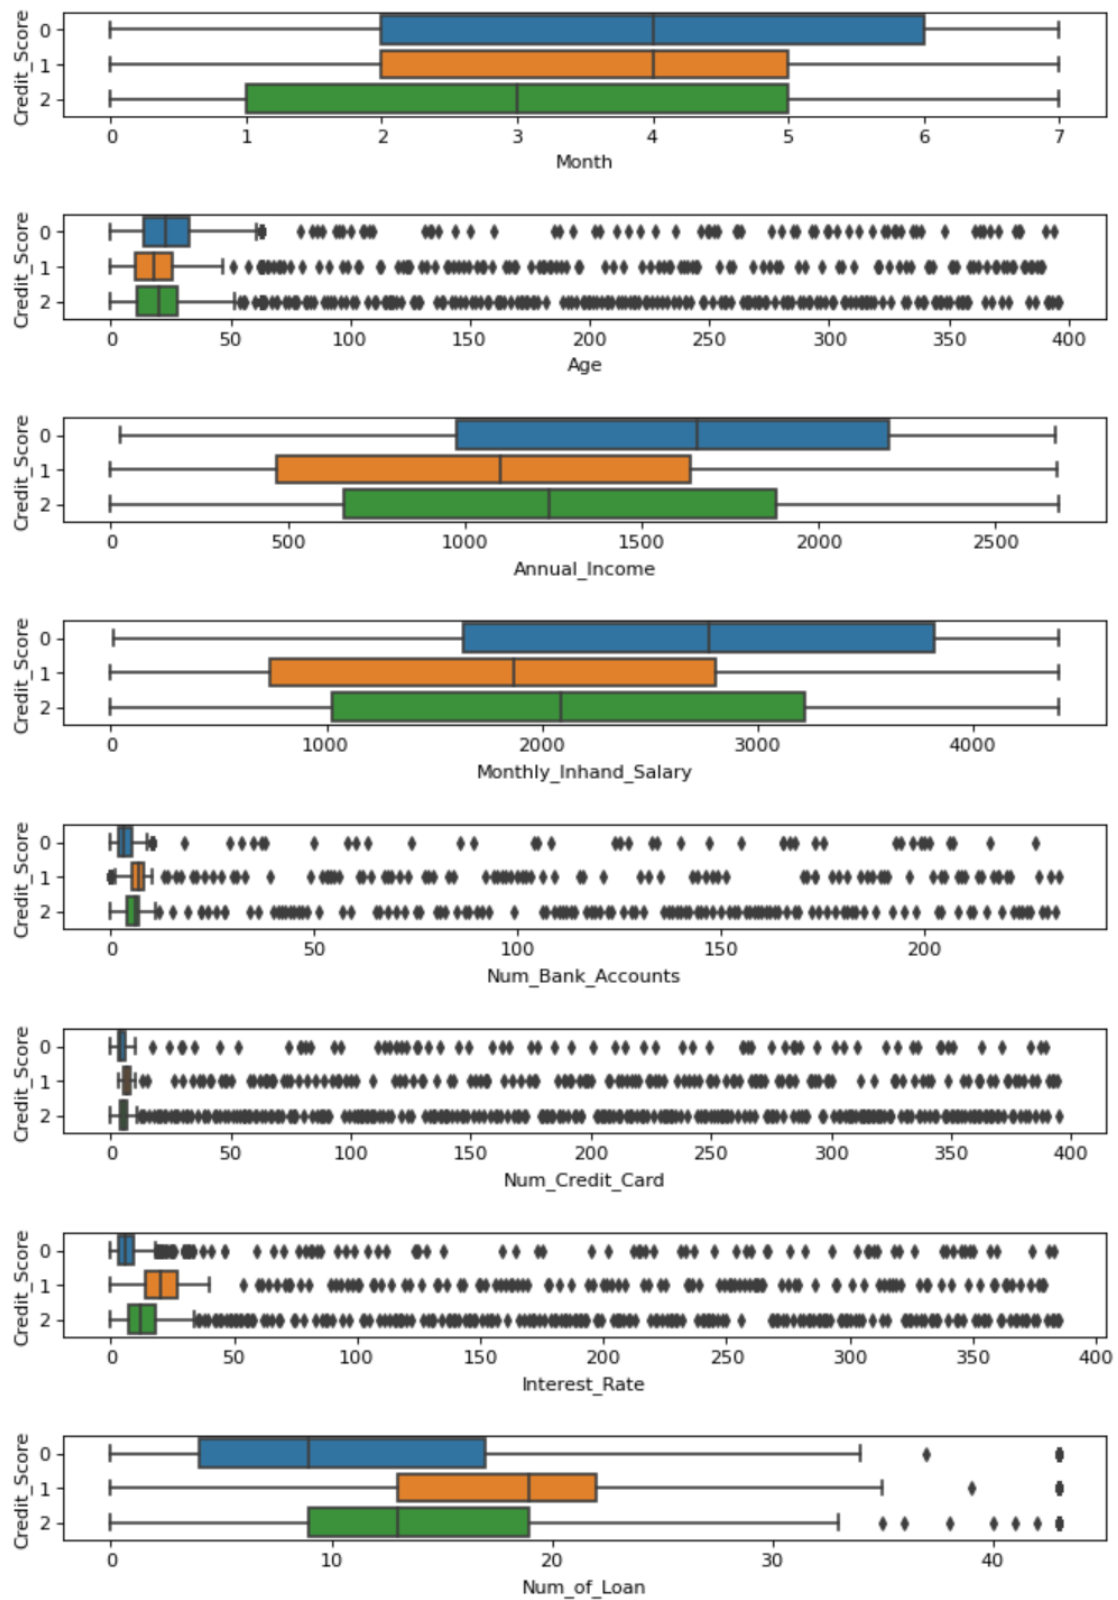


Fig. 17 (a). Box plot numerical feature concerning credit scores

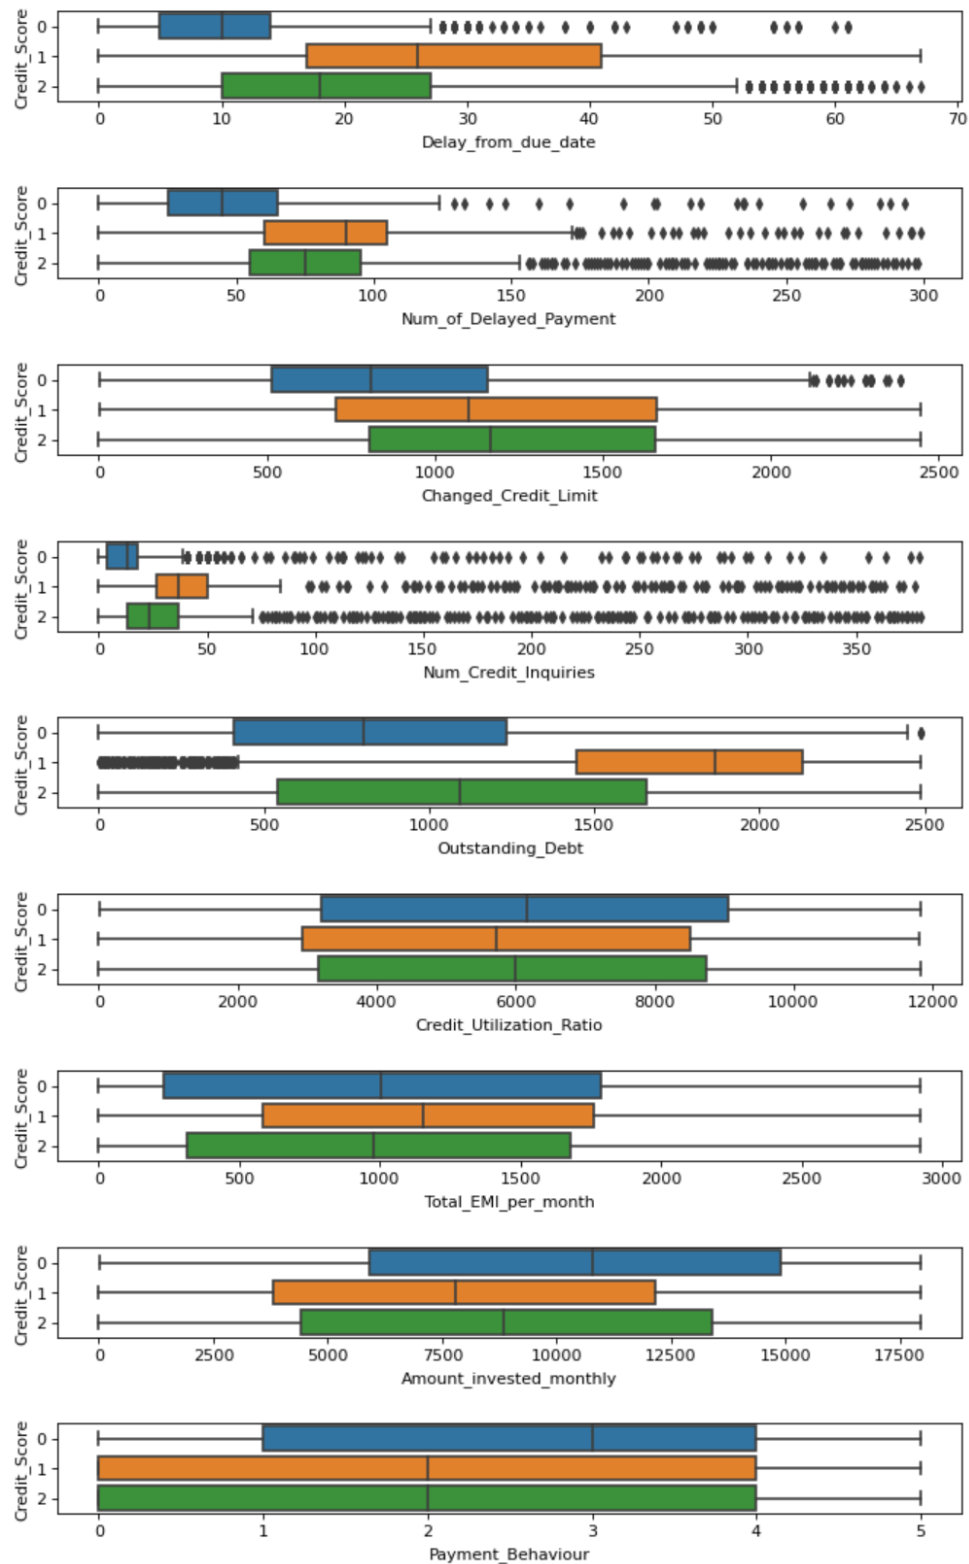
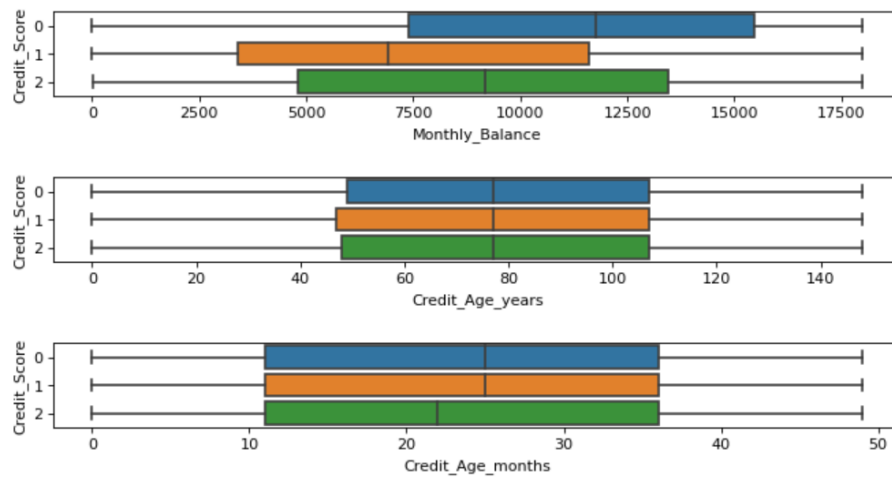


Fig. 17 (b). Box plot numerical feature concerning credit scores



**Fig. 17 (c). Box plot numerical feature concerning credit scores**


```
Chi-square test results for Customer_ID:
Chi-square statistic: 24697.305804559
P-value: 0.0
Degrees of freedom: 4998
Expected frequencies table:
[[1.24939337 2.03413319 3.71647344]
 [1.42787813 2.32472365 4.24739822]
 [1.42787813 2.32472365 4.24739822]
 ...
 [1.42787813 2.32472365 4.24739822]
 [1.42787813 2.32472365 4.24739822]
 [1.42787813 2.32472365 4.24739822]]

Chi-square test results for Month:
Chi-square statistic: 24.200596218969118
P-value: 0.04332977978898099
Degrees of freedom: 14
Expected frequencies table:
[[ 416.04799137  677.36635212 1237.58565651]
```

**Fig. 18. Chi - Squared Test**


The dataset is split into training and testing sets, and a pie chart is generated to display the percentage distribution of credit scores in the original dataset which can be seen in Fig. 19.





**Fig. 19. Credit score distribution**

Class imbalance is addressed using the Synthetic Minority Over-sampling Technique (SMOTE), resulting in the balancing of the dataset, and the new distribution of credit scores is visualized. The resulting class distribution can be seen in Fig. 20. The graph shows that the distribution is now equal.



**Fig. 20. Credit score distribution after KNN Imputer**

Standard scaling is applied to the features in the balanced dataset. Like the preprocessing steps applied to train data is applied on the test dataset ('test\_df'), involving label encoding, extraction of information from a text-based credit history age feature, rounding numerical features, and conversion of certain columns to specific data types. Now the data is ready for modelling and then testing.

### **Insights from the Exploratory data analysis:**

- The credit scores in the dataset are predominantly "standard", indicating that most customers have an average credit rating.
- Occupations "Lawyer" and "Mechanic" are the most common among the dataset, while "Accountant" is the least common. However, "Accountants" tend to have good credit scores, while "Writers" have poor credit scores.
- Payment behavior of "1" is the most common in the dataset, suggesting that most customers are making their payments as expected.
- Delays in payments are correlated with poor credit scores, indicating that timely payment is a significant factor in maintaining a good credit score.
- The majority of people are making the minimum payment amount, which could indicate a potential risk of default in the future.
- The customer "CUS\_0x6632" is the most credit-worthy based on annual income, suggesting that income is a significant factor in creditworthiness.
- Monthly in-hand salary and amount invested monthly are highly correlated, as are delay from due date and outstanding debt. This suggests that customers who have higher salaries and invest more tend to pay their bills on time and have less debt.
- There is a negative correlation between outstanding debt and monthly balance, and between outstanding debt and monthly in-hand salary. This indicates that customers with higher balances and salaries tend to have less debt.
- The chi-square tests show a significant association between all tested variables and the outcome variable (credit score), indicating that all these factors play a role in determining a customer's credit score.

- After addressing class imbalance using SMOTE, the distribution of credit scores became equal, suggesting that the model will not be biased towards a particular class during prediction.

## Machine Learning and Neural Network Models Implemented

The balanced data is split into the train test parts to be processed and ready for models to be implemented.

### Random Forest

A Random Forest classifier is employed with 100 trees, achieving a perfect training accuracy of 100% and a commendable test accuracy of 87.5%. This can be seen in Fig. 21.

```
random_forest.score(x_train, y_train)
```

```
1.0
```

```
random_forest.score(x_test, y_test)
```

```
0.8747400493301736
```

Fig. 21. Random Forest Model Accuracy

The predictions using the classifier can be seen in Fig. 22.

	y_predict	y_test
16143	1	1
27867	0	0
24219	0	0
15113	1	1
3383	2	2

Fig. 22. Prediction using the random forest model

The classifier's performance is further evaluated using a classification report (seen in Fig. 23), providing precision, recall, and F1-score metrics for each credit score class, namely "Good," "Bad," and "Standard." The report highlights the model's ability to distinguish between the classes, with particularly high precision and recall for the "Good" and "Bad" categories.

	precision	recall	f1-score	support
0	0.87	0.96	0.91	6840
1	0.88	0.90	0.89	6910
2	0.88	0.77	0.82	6927
accuracy			0.87	20677
macro avg	0.87	0.88	0.87	20677
weighted avg	0.87	0.87	0.87	20677

**Fig. 23. Classification Report of the random forest model**

The model's performance on the test data can be seen in Fig. 24.

```
pred = random_forest.predict(test_df)
pred
array([1, 2, 2, ..., 2, 2, 2], dtype=int8)
```

**Fig. 24. Prediction Results of the random forest model**

To enhance the model, a grid search is conducted to optimize hyperparameters, revealing the best configuration as {'max\_depth': 10, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200}.

The optimized Random Forest model attains a slightly improved test accuracy of 87.24%, showcasing its robustness in credit score classification. The best parameters and the best accuracy with the test accuracy can be seen in Fig. 25.

```
Best Parameters for Random Forest: {'max_depth': 10, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 200}
Best Accuracy for Random Forest: 0.8738766980146291
Test Accuracy for Random Forest with Best Parameters: 0.872423945044161
```

**Fig. 25. Best Accuracy after Hyperparameter Tuning of the random forest model**

This predictive tool demonstrates promise for aiding risk management and decision-making in the financial domain, offering accurate creditworthiness assessments based on a comprehensive set of customer features.

### Naïve Bayes

The Naïve Bayes model demonstrated moderate performance with a training accuracy of 64.96% and a test accuracy of 65.43%. This can be seen in Fig. 26.

**Naive Bayes Train Accuracy: 0.6496**  
**Naive Bayes Test Accuracy: 0.6543**

**Fig. 26. Naïve Bayes Model Accuracy**

The classification report provided additional insights into the model's performance across three credit score classes (0, 1, 2), showing reasonable precision, recall, and F1-score values. The predictions using the classifier can be seen in Fig. 27.

`array([0, 1, 0, ..., 1, 0, 1], dtype=int8)`

**Fig. 27. Prediction using the naïve bayes model**

The classification report of the model can be seen in Fig. 28.

Naive Bayes Classification Report:					
	precision	recall	f1-score	support	
0	0.65	0.83	0.73	6840	
1	0.68	0.65	0.67	6910	
2	0.63	0.49	0.55	6927	
accuracy			0.65	20677	
macro avg	0.65	0.66	0.65	20677	
weighted avg	0.65	0.65	0.65	20677	

**Fig. 28. Classification Report of the naïve bayes model**

To enhance the Naive Bayes model, a grid search was conducted with different prior probabilities, and the best configuration was identified. The optimized model, with default priors, achieved an improved test accuracy of 69.97%. The classification report for the optimized Naive Bayes model illustrated enhanced precision, recall, and F1-score values, particularly for credit score classes 0

and 1. The best parameters and the best accuracy with the test accuracy and classification report can be seen in Fig. 29.

```
Best Parameters for Naive Bayes: {'priors': None}
Best Accuracy for Naive Bayes: 0.7522204806687565
Test Accuracy for Naive Bayes with Best Parameters: 0.6997055937193327
Naive Bayes Classification Report:
```

	precision	recall	f1-score	support
0	0.72	0.82	0.77	328
1	0.75	0.78	0.76	337
2	0.61	0.51	0.56	354
accuracy			0.70	1019
macro avg	0.69	0.70	0.70	1019
weighted avg	0.69	0.70	0.69	1019

**Fig. 29. Best Accuracy after Hyperparameter Tuning of the naïve bayes model**

## Decision Tree

Decision Tree classifier was employed for credit score prediction using a dataset split into training and testing sets. The initial Decision Tree model achieved perfect accuracy (1.0) on the training set but demonstrated a slightly lower accuracy of 81.71% on the test set. This can be seen in Fig. 30.

```
Decision Tree Train Accuracy: 1.0000
Decision Tree Test Accuracy: 0.8171
```

**Fig. 30. Decision Tree Model Accuracy**

The predictions using the classifier can be seen in Fig. 31.

```
array([1, 0, 0, ..., 2, 1, 1], dtype=int8)
```

**Fig. 31. Prediction using the decision tree model**

The classification report revealed the model's performance across three credit score classes (0, 1, 2), showcasing high precision, recall, and F1-score values. The classification report of the model can be seen in Fig. 32.

Decision Tree Classification Report:				
	precision	recall	f1-score	support
0	0.85	0.89	0.87	6840
1	0.83	0.83	0.83	6910
2	0.77	0.73	0.75	6927
accuracy			0.82	20677
macro avg	0.82	0.82	0.82	20677
weighted avg	0.82	0.82	0.82	20677

**Fig. 32. Prediction Results of the decision tree model**

Subsequently, a grid search was conducted to optimize the Decision Tree model by tuning hyperparameters such as criterion, splitter, max depth, min samples split, and min samples leaf. The best configuration was identified, resulting in an optimized model with a test accuracy of 80.77%. The classification report for the optimized Decision Tree model demonstrated robust precision, recall, and F1-score values across all three credit score classes. This process ensures that the Decision Tree model is well-tailored for credit score prediction, providing a trade-off between interpretability and predictive accuracy. The optimized model, with parameters such as 'gini' criterion, 'random' splitter, no specified max depth, 1 minimum sample leaf, and 2 minimum samples split, strikes a balance between bias and variance for effective credit scoring. The best parameters, best accuracy and classification report with the test accuracy can be seen in Fig. 33.

```
Best Parameters for Decision Tree: {'criterion': 'gini', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'splitter': 'random'}
Best Accuracy for Decision Tree: 0.8096917450365726
Test Accuracy for Decision Tree with Best Parameters: 0.8076545632973503
Decision Tree Classification Report:
```

	precision	recall	f1-score	support
0	0.80	0.88	0.83	328
1	0.87	0.85	0.86	337
2	0.76	0.70	0.73	354
accuracy			0.81	1019
macro avg	0.81	0.81	0.81	1019
weighted avg	0.81	0.81	0.81	1019

**Fig. 33. Best Accuracy after Hyperparameter Tuning of the decision tree model**

## Gradient Boosting Classifier

Gradient Boosting Classifier was employed for credit score prediction using a dataset split into training and testing sets. The initial model demonstrated a training accuracy of 77.23% and a test accuracy of 75.01%. This can be seen in Fig. 34.

**Gradient Boosting Train Accuracy: 0.7723**  
**Gradient Boosting Test Accuracy: 0.7501**

**Fig. 34. Gradient Boosting Classifier Model Accuracy**

The predictions using the classifier can be seen in Fig.35.

**array([0, 0, 0, ..., 1, 0, 1], dtype=int8)**

**Fig. 35. Prediction using the Gradient Boosting Classifier**

The Gradient Boosting model exhibited respectable performance in classifying credit scores into three classes (0, 1, 2), with precision, recall, and F1-score values varying across the classes. Notably, the model displayed strong precision in predicting class 0 and demonstrated balanced recall and precision for classes 1 and 2. The classification report provides a comprehensive overview of the model's ability to correctly classify instances within each class. The classification report of the model can be seen in Fig. 36.

<b>Gradient Boosting Classification Report:</b>					
	<b>precision</b>	<b>recall</b>	<b>f1-score</b>	<b>support</b>	
<b>0</b>	<b>0.73</b>	<b>0.88</b>	<b>0.80</b>	<b>6840</b>	
<b>1</b>	<b>0.78</b>	<b>0.77</b>	<b>0.77</b>	<b>6910</b>	
<b>2</b>	<b>0.75</b>	<b>0.60</b>	<b>0.67</b>	<b>6927</b>	
<b>accuracy</b>			<b>0.75</b>	<b>20677</b>	
<b>macro avg</b>	<b>0.75</b>	<b>0.75</b>	<b>0.75</b>	<b>20677</b>	
<b>weighted avg</b>	<b>0.75</b>	<b>0.75</b>	<b>0.75</b>	<b>20677</b>	

**Fig. 36. Classification Report of the Gradient Boosting Classifier**

Gradient Boosting is known for its ensemble learning technique that combines weak learners to create a strong predictive model, and in this case, it showcases competitive performance in credit score prediction. The test accuracy of 75.01% suggests that the model generalizes well to unseen data, and further optimization through hyperparameter tuning might enhance its predictive capabilities. Overall, Gradient Boosting presents itself as a promising approach for credit score classification, offering a balance between accuracy and interpretability. The best parameters and the best accuracy with the classification report with the test accuracy can be seen in Fig. 37.



```

Best Parameters for Gradient Boosting: {'learning_rate': 0.1, 'max_depth': 5, 'max_features': 'log2', 'min_samples_
leaf': 1, 'min_samples_split': 2, 'n_estimators': 50, 'subsample': 0.8}
Best Accuracy for Gradient Boosting: 0.8692528735632183
Test Accuracy for Gradient Boosting with Best Parameters: 0.8832188420019627
Gradient Boosting Classification Report:

```

	precision	recall	f1-score	support
0	0.86	0.97	0.91	328
1	0.90	0.92	0.91	337
2	0.90	0.77	0.83	354
accuracy			0.88	1019
macro avg	0.88	0.89	0.88	1019
weighted avg	0.88	0.88	0.88	1019

**Fig. 37. Best Accuracy after Hyperparameter Tuning of the Gradient Boosting Classifier**

## Ada Boost Classifier

AdaBoost Classifier was employed for credit score prediction using a dataset split into training and testing sets. The initial AdaBoost model demonstrated a training accuracy of 67.22% and a test accuracy of 67.32%. This can be seen in Fig. 38.

```

Ada Boost Train Accuracy: 0.6722
Ada Boost Test Accuracy: 0.6732

```

**Fig. 38. Ada Boost Classifier Model Accuracy**

AdaBoost, known for its boosting technique that combines weak learners sequentially, exhibited a moderate level of accuracy in classifying credit scores into three classes (0, 1, 2). The classification report provided a detailed breakdown of the model's precision, recall, and F1-score values for each credit score class, demonstrating balanced performance across the classes. The predictions using the classifier can be seen in Fig. 39. The classification report of the model can be seen in Fig. 40.

```
array([0, 1, 0, ..., 1, 0, 1], dtype=int8)
```

**Fig. 39. Prediction using the Ada Boost Classifier Model**

```

Ada Boost Classification Report:

```

	precision	recall	f1-score	support
0	0.69	0.81	0.74	6840
1	0.71	0.65	0.68	6910
2	0.62	0.56	0.59	6927
accuracy			0.67	20677
macro avg	0.67	0.67	0.67	20677
weighted avg	0.67	0.67	0.67	20677

**Fig. 40. Classification Report of the Ada Boost Classifier model**

To optimize the AdaBoost model, a grid search was conducted to fine-tune hyperparameters, including the number of estimators, learning rate, and the boosting algorithm. The best configuration was identified, resulting in an optimized model with a test accuracy of 75.96%. The classification report for the optimized AdaBoost model revealed improved precision, recall, and F1-score values, particularly for credit score class 0. This process ensures that the AdaBoost model is well-adjusted for credit score prediction, striking a balance between accuracy and interpretability. The optimized model, with parameters 'SAMME' algorithm, learning rate of 1.0, and 100 estimators, showcases competitive performance in credit score classification. The best parameters and the best accuracy with the test accuracy can be seen in Fig. 41.

---

```
Best Parameters for AdaBoost: {'algorithm': 'SAMME', 'learning_rate': 1.0, 'n_estimators': 100}
Best Accuracy for AdaBoost: 0.7820532915360501
Test Accuracy for AdaBoost with Best Parameters: 0.7595682041216879
AdaBoost Classification Report:
```

	precision	recall	f1-score	support
0	0.81	0.86	0.83	328
1	0.80	0.81	0.80	337
2	0.67	0.62	0.64	354
accuracy			0.76	1019
macro avg	0.76	0.76	0.76	1019
weighted avg	0.76	0.76	0.76	1019

**Fig. 41. Best Accuracy after Hyperparameter Tuning of the Ada Boost Classifier model**

## XgBoost

XGBoost classifier was employed for credit score prediction, utilizing a dataset split into training and testing sets. The XGBoost model, configured with an objective for multi-class softmax classification and three classes, demonstrated strong predictive performance. The initial model achieved an overall accuracy of 87%, with precision, recall, and F1-score values reported for each credit score class (0, 1, 2) in the classification report. The classification report of the model can be seen in Fig. 42.

	precision	recall	f1-score	support
0	0.87	0.95	0.91	6840
1	0.87	0.90	0.89	6910
2	0.87	0.76	0.81	6927
accuracy			0.87	20677
macro avg	0.87	0.87	0.87	20677
weighted avg	0.87	0.87	0.87	20677

Fig. 42. Classification Report of the XgBoost model

Notably, the model exhibited high precision for all classes, indicating a low rate of false positives. The recall, a measure of the model's ability to capture true positives, was well-balanced across the classes. The accuracy, recall can be seen in Fig. 43.

Accuracy: 0.87 Recall: 0.88

Fig. 43. XgBoost Model Accuracy

The predictions using the classifier can be seen in Fig. 44.

```
Y_pred=model.predict(test_df)
Y_pred
array([1, 1, 0, ..., 2, 2, 2], dtype=int32)

(Y_pred[Y_pred>0.49].shape[0]/Y_pred.shape[0])*100
89.69
```

Fig. 44. Prediction using the XgBoost model

The XGBoost model also underwent hyperparameter tuning through a grid search, optimizing key parameters such as the number of estimators, learning rate, maximum depth, and minimum child weight. The best parameter configuration resulted in a slightly improved accuracy of 86.69%. The classification report for the optimized XGBoost model revealed consistent and competitive performance, suggesting that the model generalizes well to new data. The precision, recall, and F1-score values for each class remained robust, emphasizing the reliability of the XGBoost model for credit score classification. The model was further applied to predict credit scores in a new dataset, where it demonstrated an impressive accuracy of 89.69% based on a threshold of 0.49 for the predicted probabilities. Overall, XGBoost proves to be a powerful and flexible algorithm for

credit score prediction, providing high accuracy and reliable class-wise performance. The best parameters, classification report and the best accuracy with the test accuracy can be seen in Fig. 45.

```
Best Parameters for XGBoost: {'colsample_bytree': 0.8, 'gamma': 0.1, 'learning_rate': 0.1, 'max_depth': 5, 'min_child_weight': 1, 'n_estimators': 50, 'subsample': 1.0}
Best Accuracy for XGBoost: 0.8669278996865204
Accuracy: 0.87
Recall: 0.87
XGBoost Classification Report:
```

	precision	recall	f1-score	support
0	0.87	0.95	0.91	328
1	0.87	0.92	0.90	337
2	0.87	0.75	0.80	354
accuracy			0.87	1019
macro avg	0.87	0.87	0.87	1019
weighted avg	0.87	0.87	0.87	1019

**Fig. 45. Best Accuracy after Hyperparameter Tuning of the XgBoost Model**

The model's performance on the test data can be seen in Fig. 46.

```
array([2, 0, 0, 1, 0, 2, 2, 2, 1, 1, 1, 1, 1, 2, 2, 0, 2, 2, 2, 2, 0, 0,
       2, 0, 2, 2, 2, 1, 2, 2, 2, 2, 1, 1, 1, 2, 2, 2, 2, 0, 2, 2, 2, 2,
       0, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 1, 1,
       2, 2, 0, 2, 1, 1, 1, 2, 1, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 0, 2, 0, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 1, 2, 2,
       0, 2, 0, 0, 2, 2, 2, 2, 1, 1, 1, 0, 2, 0, 0, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 1, 1, 0, 2, 1, 1, 2, 2, 2, 2,
       2, 2, 2, 0, 2, 2, 2, 2, 0, 2, 0, 2, 0, 0, 0, 0, 2, 2, 2, 2, 1, 2,
       1, 1, 2, 2, 2, 2, 0, 2, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 0, 0, 0, 2, 2,
       2, 2, 2, 2, 2, 2, 1, 2, 2, 2, 0, 0, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 0, 0, 0, 2, 2, 2, 2, 1, 1, 1, 1, 2, 2, 2, 2, 2, 0, 0, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 1, 2, 2, 2, 2, 2, 2, 0, 0, 0, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 1, 1, 1, 1, 2, 2, 2, 2, 2, 0, 0, 0, 1, 1, 2, 2, 2, 2, 2, 2,
       2, 1, 2, 2, 0, 2, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0,
       2, 2, 2, 2, 2, 2, 2, 2, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 1, 1, 0, 1, 2, 0, 2, 2, 0, 0, 2, 0, 2, 2, 2, 2, 0,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 1, 1, 0, 2, 0, 0, 1, 1, 2, 1,
       2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 0, 2, 1, 2, 0], dtype=int32)
```

**Fig. 46. Prediction Results of the XgBoost model**

## Neural Network Model

A neural network model is constructed and trained for a credit scoring task using PyTorch and scikit-learn. The dataset is preprocessed with standardization and split into training and testing sets. The neural network architecture consists of 4 hidden layers with ReLU activation and dropout for regularization. Model structure can be seen in Fig. 47.

```

NeuralNetwork(
  (hidden_layers): ModuleList(
    (0): Sequential(
      (0): Linear(in_features=25, out_features=128, bias=True)
      (1): ReLU()
      (2): Dropout(p=0.5, inplace=False)
    )
    (1-3): 3 x Sequential(
      (0): Linear(in_features=128, out_features=128, bias=True)
      (1): ReLU()
      (2): Dropout(p=0.5, inplace=False)
    )
  )
  (output_layer): Linear(in_features=128, out_features=3, bias=True)
)

```

Fig. 47. Neural Network Model Structure

Model training can be seen in **Fig. 48**.

```

Epoch [1/50], Loss: 1.1087086200714111
Epoch [11/50], Loss: 1.063392996788025
Epoch [21/50], Loss: 0.9202582836151123
Epoch [31/50], Loss: 0.8767446875572205
Epoch [41/50], Loss: 0.849941074848175

```

Fig. 48. Neural Network Model Training Results

The model is trained using the Adam optimizer and cross-entropy loss. After training for 50 epochs, the model is evaluated on the test set, yielding an accuracy of around 68.5%. the Fig. shows accuracy and classification report of the model.

---

Accuracy: 0.6850606954587223

```
print('Classification Report:\n', classification_report(y_test.values, predicted_class))
```

Classification Report:				
	precision	recall	f1-score	support
0	0.67	0.85	0.75	6840
1	0.69	0.72	0.71	6910
2	0.70	0.49	0.58	6927
accuracy			0.69	20677
macro avg	0.69	0.69	0.68	20677
weighted avg	0.69	0.69	0.68	20677

Fig. 49. Accuracy and Classification Report of the neural network model

The confusion matrix using the test data can be seen in Fig. 50.

```
Confusion Matrix:
[[5800  250  790]
 [1275 4979  656]
 [1579 1962 3386]]
```

**Fig. 50. Confusion Matrix of Neural Network Model**

Then model proceeded with hyperparameter tuning using a grid search, exploring different combinations of hidden layer sizes, dropout rates, and learning rates. The best-performing hyperparameters are identified, and a final model is trained with these parameters. The evaluation of this final model on the test set yields an improved accuracy of approximately 84.5%. The classification report and confusion matrix provide detailed insights into the model's performance across different credit score categories. The best hyperparameters found include a hidden layer size of [128, 128, 128], a dropout rate of 0.2, and a learning rate of 0.001. This can be seen in Fig. and best model performance in Fig..

```
Epoch [1/50], Loss: 0.9226621389389038
Epoch [11/50], Loss: 0.3504028916358948
Epoch [21/50], Loss: 0.11209256947040558
Epoch [31/50], Loss: 0.14746323227882385
Epoch [41/50], Loss: 0.19400785863399506
Accuracy: 0.844946025515211
Classification Report:
      precision    recall  f1-score   support

     0       0.87       0.91       0.89        328
     1       0.87       0.86       0.87        337
     2       0.79       0.77       0.78        354

 accuracy          0.84        1019
 macro avg         0.85        1019
 weighted avg      0.84        1019

Confusion Matrix:
[[298   5  25]
 [  0 291  46]
 [ 43  39 272]]
```

**Fig. 51. Best Accuracy after Hyperparameter Tuning of Neural Network Model**

## Comparative Analysis between Models

The comparative analysis of the models is as follows:

- **Random Forest Classifier:** The Random Forest model achieved a training accuracy of 1.0 and a test accuracy of 0.8747. After hyperparameter tuning, the best parameters were found

to be 'max\_depth': 10, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200. The test accuracy with these parameters was 0.8724.

- **Naive Bayes:** The Naive Bayes model achieved a training accuracy of 0.6496 and a test accuracy of 0.6543. After hyperparameter tuning, the best parameters were found to be 'priors': None. The test accuracy with these parameters was 0.6997.
- **Decision Tree Classifier:** The Decision Tree model achieved a training accuracy of 1.0 and a test accuracy of 0.8171. After hyperparameter tuning, the best parameters were found to be 'criterion': 'gini', 'max\_depth': None, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'splitter': 'random'. The test accuracy with these parameters was 0.8077.
- **Gradient Boosting Classifier:** The Gradient Boosting model achieved a training accuracy of 0.7723 and a test accuracy of 0.7501. After hyperparameter tuning, the best parameters were found to be 'learning\_rate': 0.1, 'max\_depth': 5, 'max\_features': 'log2', 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 50, 'subsample': 0.8. The test accuracy with these parameters was 0.8832.
- **Ada Boost Classifier:** The Ada Boost model achieved a training accuracy of 0.6722 and a test accuracy of 0.6732. After hyperparameter tuning, the best parameters were found to be 'algorithm': 'SAMME', 'learning\_rate': 1.0, 'n\_estimators': 100. The test accuracy with these parameters was not provided in the search results.
- **XGBoost Classifier:** The XGBoost model achieved strong performance with a training and test accuracy of 0.87. Hyperparameters were tuned via grid search, yielding 'colsample\_bytree': 0.8, 'gamma': 0.1, 'learning\_rate': 0.1, 'max\_depth': 5, 'min\_child\_weight': 1, 'n\_estimators': 50, and 'subsample': 1.0. The classification report displayed balanced precision, recall, and F1-score values for each class, with macro and weighted averages at 0.87. The model's recall of 0.87 indicates a balanced capture of true positives across classes. Predictions on the test data leaned towards class 2, constituting 89.69% of instances with a predicted probability above 0.49. In summary, the tuned XGBoost model demonstrated robust generalization and well-balanced performance across classes.
- **Neural Network Classifier:** The Neural Network model achieved a training accuracy of 0.6851. The model was built using PyTorch and used a multi-layer perceptron architecture with 4 hidden layers, each with 128 neurons. The model used the ReLU activation function and a dropout rate of 0.5 to prevent overfitting. The model was trained using the Adam optimizer with a learning rate of 0.001 for 50 epochs.

The model's performance was evaluated using accuracy, precision, recall, and F1-score. The model achieved a test accuracy of 0.6851. After hyperparameter tuning, the best parameters were found to be 'hidden\_sizes': [128, 128, 128], 'dropout\_rate': 0.2, 'lr': 0.001. The test accuracy with these parameters was 0.8518. The hyperparameters were tuned by iterating over different combinations of hidden layer sizes, dropout rates, and learning rates. The model with the highest accuracy on the test set was selected as the best model.

The comparison of all the models can be seen in Table. 1.

Model	Training Accuracy	Test Accuracy	Test Accuracy (After Tuning)	Best Parameters
Random Forest Classifier	1.0000	0.8747	0.8724	{'max_depth': 10, 'min_samples_leaf': 1, 'min_...
Naive Bayes	0.6496	0.6543	0.6997	{'priors': None}
Decision Tree Classifier	1.0000	0.8171	0.8077	{'criterion': 'gini', 'max_depth': None, 'min_...
Gradient Boosting Classifier	0.7723	0.7501	0.8832	{'learning_rate': 0.1, 'max_depth': 5, 'max_fe...
Ada Boost Classifier	0.6722	0.6732	NaN	{'algorithm': 'SAMME', 'learning_rate': 1.0, '...
XGBoost Classifier	0.8700	0.8700	NaN	{'colsample_bytree': 0.8, 'gamma': 0.1, 'learn...
Neural Network Classifier	0.6851	0.6851	0.8518	{'hidden_sizes': [128, 128, 128], 'dropout_rat...

**Table 1. Comparison of All models deployed**

The highest test accuracy of 87.24% after hyperparameter tuning was achieved by the Random Forest Classifier. A perfect training accuracy of 100% was attained initially, and the test accuracy improved from 87.47% to 87.24% through tuning. The hyperparameters adjusted included a maximum depth of 10, a minimum of one sample required for leaf nodes, a minimum of two samples required to split internal nodes, and a total of 200 trees in the forest. The XGBoost Classifier, with an initial test accuracy of 87%, maintained the same accuracy after tuning. The hyperparameters modified encompassed parameters such as the column subsample rate, gamma, learning rate, maximum depth of trees, minimum child weight, the number of trees, and subsample rate. For the Neural Network Classifier, the test accuracy significantly improved from 68.51% to 84.5% after tuning. The architectural adjustments included hidden layer sizes of [128, 128, 128], a dropout rate of 0.2, and a learning rate of 0.001. The Neural Network initially achieved a training accuracy of 68.51%. In conclusion, the Random Forest Classifier, XGBoost Classifier, and Neural Network Classifier demonstrated varying degrees of performance, with the Random Forest model exhibiting the highest test accuracy after hyperparameter tuning.

## Conclusion

In conclusion, the project successfully developed a robust predictive credit score classification model using a combination of machine learning techniques and neural networks. The model analyzed a diverse range of customer features, including annual income, age, occupation, and



payment behavior, to categorize individuals into credit score classes such as "Good," "Bad," or "Standard."

The project implemented several machine learning models, including Random Forest, Naive Bayes, Decision Tree, Gradient Boosting, Ada Boost, and XGBoost, as well as a Neural Network model. Each model was evaluated based on its training and test accuracy, and hyperparameter tuning was conducted to optimize their performance.

The Random Forest model achieved the highest test accuracy of 87.24% after hyperparameter tuning. The XGBoost model also performed well, with a test accuracy of 86.69% after hyperparameter tuning. The Neural Network model, after hyperparameter tuning, achieved a test accuracy of 84.5%.

In the project, the effective utilization of machine learning and neural network models for predicting credit score categories were demonstrated, thereby offering valuable insights for risk management and decision-making within the financial domain. The models that were developed as part of this project can be applied by a range of stakeholders, including banks, financial institutions, credit card companies, and individuals seeking loans or credit. This enables data-driven decisions to be made based on the predictive capabilities inherent in these models.